

TrustEpisode: Auditable Evidence Contracts for AI-Assisted EDR/NDR Investigation

Anonymous Author(s)

Abstract

AI-assisted security investigation turns sparse tool outputs into confident narratives, yet arrival order, collector failure, and late evidence can silently alter the facts exposed to an analyst or agent. We present TRUSTEPISODE, a deterministic, auditable EDR/NDR episode-reconstruction contract. Episode Synthesis alone owns membership; revisions and terminal audit outcomes are append-only; source health differs from “no detection;” unsupported probabilities remain null; and downstream AI reads versioned evidence without authority over case facts.

On 70 sealed live CALDERA M1–M7 malicious runs (40 standalone benign controls are not detection trials), sliding sessionization raises mean coverage from 0.807 to 0.821. Fair sliding baselines show no superiority: EB4 ties EB2 on all action-coverage pairs (0/70/0 W/T/L) and trails EB1 on one concurrent-benign M7 run (0/69/1; mean $\Delta = -0.014$, 95% CI $[-0.043, 0.000]$). Canonical replay is byte-identical on 60 M1–M6 runs, whereas raw arrival order is not robust without re-canonicalization. 18/18 audit fixtures cover seven terminal failure codes; an A–G taxonomy localizes M2’s repeatable gap to class D (token mapping failure); and a deployed case trace preserves one Suricata and two Wazuh events with ATT&CK references.

All 8/8 synthetic boundary cards behave as specified: the valid control is accepted, six invalid or unsealed mutations are rejected, and a resealed allowed-field mutation is accepted, showing digest integrity still requires external writer authorization. A once-run held-out 20-card Qwen2.5-7B pilot yields 9/20 fully exact responses; locked same-card replication on deployment-default Gemma3-1B yields 0/20. These audited configurations support no general LLM-accuracy or fitted-calibration claim; typed evidence instead makes their errors reviewable.

CCS Concepts

• **Security and privacy** → **Intrusion/anomaly detection and malware mitigation**; • **Computing methodologies** → *Artificial intelligence*.

Keywords

AI-assisted security operations, attack episode reconstruction, EDR, NDR, replayable audit, evidence provenance

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

AI-Sec 2026, The Hague, Netherlands

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2026/11
<https://doi.org/XXXXXXX.XXXXXXX>

ACM Reference Format:

Anonymous Author(s). 2026. TrustEpisode: Auditable Evidence Contracts for AI-Assisted EDR/NDR Investigation. In *Proceedings of 19th ACM Workshop on Artificial Intelligence and Security (AI-Sec 2026)*. ACM, New York, NY, USA, 11 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

AI-assisted SOC systems increasingly use language-model agents to gather tool outputs, correlate alerts, and draft incident narratives [5, 21, 25]. Their output can be no more reviewable than the evidence interface below it: if arrival order changes episode membership, collector failure is encoded as benign silence, or a late event rewrites history, a fluent explanation may cite a case state that cannot be reproduced.

Security operations and digital-forensic investigations therefore need more than alert clustering. Multi-step attacks appear as dispersed EDR and NDR observations that become material when linked by time, entity, session, and process relations [1, 11, 15]. Investigators need a contract that fixes who owns episode membership, how late evidence is admitted, how collector health differs from “no detection,” and how every revision leaves an independently verifiable terminal outcome [3, 4, 16].

Prior alert-correlation and provenance systems motivate multi-stage context using rules, graphs, and learned representations [6, 8–10, 12, 20]. Those lines do not remove the need for fixed ownership of membership, immutable revision history, and replayable audit semantics under incomplete observability. This submission evaluates TRUSTEPISODE as a *deterministic, auditable reconstruction contract* for EDR/NDR evidence and as a typed, read-only boundary for downstream analyst or AI reasoning. The downstream model is deliberately outside the trusted computing base: it may summarize versioned AuditRecords but cannot own membership, infer source health, overwrite history, or convert unsupported evidence into a numeric probability. A compact scorer and calibrator remain optional schema interfaces and are not empirically validated here. Thus, our AI-security claim concerns *evidence truthfulness and containment*, not broad LLM effectiveness. We additionally report a small, held-out Qwen pilot and a post-result locked replication on the deployment-default Gemma model to measure extraction, citation, and abstention over the typed record. They are boundary stress tests rather than a preregistered model comparison.

Research questions.

- RQ1** Can TrustEpisode reproduce byte-identical membership, revision history, late-evidence records, and digests on 60 sealed malicious runs?
- RQ2** Do open, finalized, late-successor, closed, beyond-horizon, duplicate, and all seven AssemblyFailureRecord codes yield explicit immutable audit outcomes without rewriting earlier history?

- RQ3** When a required action is absent from a reconstructed episode, can the pipeline localize loss across the closed A–G taxonomy, and is that taxonomy instantiated on sealed live evidence?
- RQ4** Which relation, hub, merge, and late-revision mechanisms does the live M1–M7 cohort exercise, and which require targeted conformance cards?
- RQ5** What portion of observed formation gain comes from replacing fixed bins with sliding sessionization, and does relation-aware logic provide additional measurable benefit on the sealed cohort?
- RQ6** Do closed schemas and digest bindings fail closed when a downstream consumer injects asserted facts, labels, or unsupported probability; removes provenance; or tampers with revision and evidence digests; what remains possible after a schema-valid reseal; and how reliably do two local models follow the resulting read-only extraction, citation, and abstention interface?

These questions evaluate the evidence substrate. The deployed interface trace reported later checks end-to-end materialization but is not counted as an effectiveness trial, and its displayed operational confidence is not the unevaluated calibrated probability field defined by the contract.

Contributions.

- (1) a single-owner reconstruction and audit contract with immutable late-data semantics, canonical references and digests, exactly one terminal AssemblyOutcome per revision, and a read-only agent boundary;
- (2) a sealed evaluation separating evidence classes: 110 live executions, fair EB0–EB4 baselines, a 60-run determinism matrix, a complete seven-code failure matrix, contract ablations, and an A–G loss taxonomy that localizes M2 to token mapping;
- (3) downstream stress tests that reject six schema-invalid or unsealed mutations, deliberately accept a resealed allowed-field negative control to expose the writer-authorization dependency, and reveal extraction/citation/abstention failures in held-out Qwen and locked same-card Gemma audits, without claiming model superiority; and
- (4) an executable anonymous companion plus a deployed read-only trace showing source-separated Suricata/Wazuh evidence and ATT&CK support at the analyst surface.

2 Related Work

2.1 Alert Correlation and Attack Reconstruction

Classical alert correlation reconstructs multi-step scenarios from intrusion alerts by temporal and causal linking [12]. HOLMES and UNICORN recover APT context from system graphs [8, 10]; KAIROS, NODLINK, and CAPTAIN extend this line [6, 9, 20]. These systems motivate episode construction. Our research unit is instead the lifecycle contract governing mixed-EDR/NDR membership, late-data revisions, and terminal audit outcomes; detection and provenance context are complementary inputs, not effectiveness baselines.

2.2 Digital Forensic Event Reconstruction

Timeline-based event reconstruction is a core digital-forensic technique for inferring past activity from heterogeneous artifacts [1, 4]. Recent work formalizes timestamp interpretation via time anchors and clock-skew anomalies [16] and surveys challenges beyond mere data volume [1]. Process models and filesystem analysis emphasize documenting what was examined, what was unavailable, and how conclusions depend on incomplete traces [3, 11]. TrustEpisode imports these audit requirements into *online* EDR/NDR episode reconstruction: immutable revisions, explicit late stores, and digestable terminal outcomes under delayed delivery.

2.3 Reproducibility, Tool Validation, and Provenance

Forensic tool and method validation requires documented procedures, known datasets, and repeatable results before casework [2, 14, 18]. Chain-of-custody and provenance discipline demand that investigators show integrity of examined objects and distinguish observed from unavailable evidence [4, 11]. TrustEpisode’s contribution is not a higher detector F1; it is a single-owner revision contract, canonical replay boundary, and explicit terminal failure semantics that make reconstruction outcomes auditable under the same reproducibility expectations.

2.4 Incomplete Telemetry and Evidence-Loss Attribution

EDR and NDR fail differently; missing modality, parser blind spots, and collector health can erase or distort attack steps without implying that correlation “missed” an already-normalized event [17, 23]. Reconstruction contracts must therefore separate sensor observability, normalization/token mapping, formation membership, and offline attribution when diagnosing incomplete episodes.

2.5 Evidence-Grounded and Contained AI Investigation

Tool-using language-model systems interleave reasoning with external observations [25]. In security operations, CORTEX coordinates specialized agents over enterprise evidence, while Holmes packages network evidence into auditable anchors for an LLM investigator [5, 21]. These systems motivate evidence-grounded AI assistance, but grounding also needs an upstream contract: citation presence does not by itself establish faithfulness [7, 19], and agent execution requires architectural containment rather than prompt-only trust [24]. TRUSTEPISODE addresses the complementary boundary below the agent. It fixes which EDR/NDR objects exist, which producer owns them, how missingness is represented, and which immutable references a later narrative may cite. We do not claim a superior LLM or better autonomous triage.

2.6 Positioning Summary

Table 1 maps research units and objectives; it is not an effectiveness leaderboard. Our unit is *versioned evidence state*, not a score, suspicious subgraph, or narrative. The novelty claim is their contract composition: one membership writer, immutable successor/late-store semantics, source health separate from silence, one terminal

outcome per revision, stage-localized loss, and citable read-only records. Contract invariants and formation effectiveness are evaluated separately; no new detector or classifier is claimed.

2.7 Scope Distinction

TRUSTEPISODE does not claim a new graph detector, AI ranking effectiveness, attribution beyond observable EDR/NDR, enterprise impact estimation, or autonomous response. The evaluated contribution is a deterministic forensic reconstruction and audit contract for sealed EDR/NDR evidence under late arrival and incomplete observability. Its AISec relevance is architectural: AI assistance is constrained to read and cite versioned case facts rather than silently redefining them.

3 TrustEpisode Reconstruction Contract

TRUSTEPISODE converts EDR/NDR telemetry into deterministic episode revisions and version-scoped forensic audit records. Evidence Fabric and Health Monitor own events and source health; Episode Synthesis alone owns membership; optional Feature Extractor / Scorer / Calibrator fields remain schema interfaces without fitted artifacts in this submission. For each revision, the AuditRecord Assembler emits exactly one terminal AssemblyOutcome: a canonical AuditRecord or an immutable AssemblyFailureRecord. Ground truth is joined only offline. Figure 1 shows these boundaries.

3.1 Investigation and Threat Model

Assumptions. Raw pointers and the sealed evidence store are integrity-checkable; Health Monitor coverage messages have authenticated sources; attackers may induce missing events, delayed delivery, and duplicate retries; attackers cannot arbitrarily forge signed health state; runtime software may fault; the offline evaluator is isolated from online formation.

Non-assumptions. We do *not* assume complete telemetry, that collector silence equals downtime, that all attacks are observable, that ground truth is available online, that model scores are reliable, or that raw arrival order is fixed.

Downstream-agent adversary. An AI assistant may be mistaken, prompt-injected, or compromised: it may omit evidence, misquote a record, or recommend an unsafe action. It is outside the trusted computing base and has no capability to write canonical evidence objects. The contract does not make its narrative true; it makes any cited case claim checkable against an immutable revision and prevents the narrative from silently becoming the forensic record.

Investigator workflow. Investigators import sealed EDR/NDR, form revisions, and inspect each terminal audit outcome. They then check coverage and provenance, resolve late evidence, run the A-G waterfall for a missing step, replay the digests, and report observed, unavailable, and unknown separately.

3.2 Evidence and Reconstruction Model

Online objects are NormalizedEvent, SourceCoverage, EpisodeRevision, LateEvidenceRecord, AuditRecord, and AssemblyFailureRecord. Offline labels never enter online schemas. Table 2 assigns every object to one producer.




Figure 1: Label-blind online reconstruction and offline evaluation boundary.

3.3 Independent Evidence Health and Observability

Event silence is not an outage. “No detection” is not “unavailable.” SourceCoverage carries authenticated health, parser validity, and an observed source mask. Assemblers consume exact coverage references; they do not infer health from empty event windows.

3.4 Deterministic Episode Formation

Formation links events using process-causal, authenticated-session, endpoint-connection, and specific-entity relations under a 300 s maximum inter-event gap and 3600 s maximum episode span. Tie-breaks are deterministic (relation precedence, relation count, Δt , start time, episode ID). Hub suppression, parent merge, and exact event-time ordering are part of the frozen policy. Formation never reads labels, receipts, or scenario metadata.

3.5 Immutable Revision Lifecycle

Watermark lag finalizes open episodes. Late evidence inside a fixed 900 s revision horizon may create an immutable successor revision; beyond-horizon or closed-lineage evidence produces a LateEvidenceRecord. Earlier revision digests are never rewritten. Deadlines are not extended by late arrivals.

3.6 Terminal Forensic Audit Contract

At revision.watermark+120 s the assembler inserts exactly one terminal outcome under (*episode_id*, *revision*): AuditRecord or AssemblyFailureRecord with a closed failure code set (schema, key, duplicate, digest, version, missing-input, timeout). Success cannot follow terminal failure. References are canonically ordered; RFC 8785/JCS serialization and SHA-256 digests bind versions. The store is append-only.

3.7 Agent-Facing Evidence Boundary

Downstream analyst views and AI assistants are consumers, never producers, of canonical case facts. They receive a revision key, terminal AssemblyOutcome, ordered evidence references, SourceCoverage references, and the independent sufficiency profile. A generated explanation or recommended action is a noncanonical view: it cannot write EpisodeRevision, change *P* or *U*, or replace an AssemblyFailureRecord with success. If required inputs are missing,

Table 1: Related-work scope map (not an effectiveness comparison).

Research line	Primary unit	Primary objective	Relation to this work
Alert correlation [12]	alerts and scenario graphs	link multi-stage intrusion alerts	motivates temporal/causal linking; different output and evaluation target
Provenance APT analysis [6, 8–10, 20]	system/provenance graphs and suspicious subgraphs	detect, rank, or contextualize attack activity	complementary signal/context producers for a versioned evidence contract
Forensic reconstruction [1, 4, 16]	artifacts, time anchors, and timelines	reconstruct prior activity and document uncertainty	motivates chain-of-custody, late-evidence, and replay semantics
Evidence-grounded agents [5, 21, 25]	tool observations, citations, and narratives	orchestrate evidence and assist investigation	downstream consumers that need stable citable object identity
TrustEPISODE (this work)	versioned EDR/NDR episodes and terminal audit outcomes	make online reconstruction deterministic, append-only, loss-localized, and citable	evaluated contract; no detector or LLM-superiority claim

Table 2: Single-owner online object contract.

Object	Sole producer	Consumers	Forbidden behavior
NormalizedEvent	Evidence Fabric	Episode Synthesis	label, scenario, or health fields
SourceCoverage	Health Monitor	Scope Resolver; Sufficiency; Assembler	infer outage from event silence
EpisodeRevision / LateEvidenceRecord	Episode Synthesis	Feature; Sufficiency; Assembler / late store	rewrite a prior revision; admit beyond-horizon into membership
FeatureVector $[D, C, m]$	Feature Extractor	Scorer; Assembler	labels or inferred health
Nullable P	Calibrator (design)	Assembler	numeric unsupported output
Structured U	Sufficiency Evaluator	Assembler	scalarization that rewrites P
AssemblyOutcome	AuditRecord Assembler	Append-only store; replay	both success and failure; update/delete

Deterministic episode and revision lifecycle




Figure 2: Episode revision lifecycle: open → finalized → late successor or closed; beyond-horizon evidence enters LateEvidenceRecord without rewriting earlier membership.

the only canonical output is the failure record; an agent is not permitted to fill the gap from conversational context. This boundary makes a later narrative falsifiable against immutable source references without treating the language model as part of the forensic root of trust. Section 4.13 tests closed-schema, support-coupling, provenance, and digest guards. Section 4.14 then treats the model as an untrusted reader and scores exact values, abstention, and pointer citations; service-identity authorization remains a deployment assumption rather than an evaluated result.

Integrity is not writer authorization. The closed schema and unkeyed RFC 8785/SHA-256 digest validate object shape and detect an unsealed change; they do not authenticate the producer. A downstream principal with canonical write permission could change an allowed field and recompute the digest. The read-only boundary therefore requires two layers: externally enforced writer authorization, plus the evaluated schema/digest contract. AI7 is an expected-accept negative control that makes this dependency explicit rather than treating a hash as a signature.

Table 3: Evidence-loss taxonomy for incomplete reconstruction.

Class	Decision rule (summary)
A_execution_failure	Action receipt/predicate shows non-success
B_sensor_observability_failure	Receipt success but raw EDR/NDR lacks identifiable activity
C_parser_or_normalization_failure	Raw present; NormalizedEvent absent or rejected token
D_token_mapping_failure	NormalizedEvent present under wrong/missing required token
E_formation_membership_failure	Correct token exists but not retained in matched episode
F_offline_attribution_failure	Episode retains evidence; matcher fails to credit token
G_unknown_or_incomplete_trace	Required supporting records missing for A–F

3.8 Evidence-Loss Taxonomy

When a required action token is absent from a reconstructed episode, diagnosis uses the following closed classes (Table 3):

Forbidden inference: treating $F1 = 1.0$ as exact-complete step recovery; treating missing tokens as formation failures without a waterfall.

3.9 Scoring and Probability Extension (Design Only)

Schema fields may retain an optional score interface: version-bound, nullable P outside support, and U that must not rewrite P . This submission reports no fitted coefficients, Brier/NLL, or held-out ranking. Scoring and calibration are not empirical claims here.

3.10 Boundary

The *evaluated* contribution is deterministic episode reconstruction and replayable audit semantics, including evidence-loss localization and a deployed read-only case-materialization trace. Manifest-scoped probability is design-only. LLM generation quality, CTI enrichment, impact scoring, autonomous response, and OT/ICS generalization are excluded.

4 Evaluation

4.1 Evidence Classes and Cohorts

- **RQ1** reproducibility — sealed offline replay on all 60 primary malicious runs.
- **RQ2** lifecycle auditability — A1–A12 synthetic audit conformance covering seven failure codes + late/beyond-horizon cards.
- **RQ3** evidence-loss localization — A–G taxonomy conformance + sealed M2 waterfall.
- **RQ4** mechanism coverage — live M1–M7 + synthetic T1–T8, with M7 as boundary stress.
- **RQ5** baseline interpretation — sealed fair EB0–EB4 comparison on M1–M7.
- **RQ6** downstream boundary and reading — synthetic AI0–AI7 canonical-record validation plus two once-run 20-card local-model grounding audits.

Evidence classes used below are disjoint: (i) primary live empirical (M1–M7 formation cells); (ii) sealed offline replay on those primary seals; (iii) synthetic conformance (A1–A12, T1–T8, AB0–AB3, AI0–AI7, A–G fixtures); (iv) two once-run synthetic held-out LLM audits; (v) unit/integration boundary fixtures; (vi) supplemental partition artifacts (L1–L4, wall-clock) that are *not* mixed into primary estimators. Synthetic cards are never mixed into primary live effectiveness. Sealed Table 5 primary cells were not overwritten.

4.2 Primary Live M1–M7 Cohort

CALDERA v5.3.0 executes frozen Bash abilities on two isolated Ubuntu 22.04 targets with an EDR wrapper and inline Suricata NDR. Online schemas reject labels and receipts. The primary sealed cohort has 70 malicious runs (M1–M7 $\times$ two modes $\times$ five) and 40 standalone benign execution controls. Because no fitted scorer or decision rule is evaluated, those controls do not define an alert false-positive rate and are not mixed into the malicious formation estimator. Fair sliding baselines and paired comparisons are computed on the M1–M7 set.

4.3 Ground Truth and Matching

Offline matching uses temporal IoU ≥ 0.10 , entity Jaccard ≥ 0.50 , overlapping required action tokens, and weight

$$w = 0.5A_{cov} + 0.3J_{entity} + 0.2IoU_{time}; \quad w \geq \theta_{match} = 0.50. \quad (1)$$

Exact-complete recovery requires $A_{cov} = 1$. Online builders never read evaluation fields.

4.4 Fair Reconstruction Baselines

- **EB0_FIXED_BIN_300s** (legacy): fixed bins; boundary reference only.
- **EB1_SLIDING_TIME**: same $W_{gap} = 300$ s, $W_{max} = 3600$ s; time only.
- **EB2_SLIDING_ENTITY**: same lifecycle; entity overlap only.
- **EB3_RELATION_NO_ADVANCED_GUARDS**: relations on; hub/merge/late off.
- **EB4_TRUSTEPISODE**: full production policy.

Primary comparison is EB1–EB4 on identical sealed streams, aliases, matching, and ground truth. Configuration digests are sealed in `results/derived/fair_baselines.json`.

4.5 Primary Formation Results and Fair Negative Result

Table 5 reports production-formation cells for the primary live M1–M7 cohort ($n = 5$ per cell). M7 is a boundary-stress family: the final action occurs at 3900 s, beyond the frozen $W_{max} = 3600$ s span, so zero coverage/F1 is contract-conformant rather than a collection failure.

Relative to legacy EB0, sliding policies raise coverage from 0.807 to 0.821 and exact-complete recovery from 0.671 to 0.714 on the sealed primary cohort. Fixed bins cut multi-minute chains at arbitrary boundaries; sliding gap/sessionization preserves those chains under the same span cap.

The fair sliding baselines are not identical. EB4 matches EB2 on all action-coverage pairs (0/70/0 W/T/L). Against EB1, action coverage and exact-complete recovery have 0/69/1 W/T/L, mean $\Delta = -0.014$, and 95% CI $[-0.043, 0.000]$; episode-F1 has mean $\Delta = -0.010$ and 95% CI $[-0.029, 0]$. The single loss is the concurrent-benign M7 run `run-ee27d58be2db032c44d0774d9358c4b1`: EB1 recovers its truth group while EB4 does not. We therefore report the EB4–EB1 result as *not statistically resolved*, not “no observed difference,” equivalence, or relation-aware superiority. EB3 shows the same one-run action-coverage pattern.

Why sliding-time already suffices on M1–M7. The sealed primary cohort does not make advanced relations the unique linkage path: EB4 and EB2 match on all 70 action-coverage pairs, and EB4 matches EB1 on 69. On the remaining concurrent-benign M7 run, EB1 and EB3 recover a truth group that EB4 and EB2 do not. Thus the advanced-guard policy incurs one observed boundary-case loss in this cohort rather than an end-to-end gain. End-to-end F1 alone can misleadingly suggest that complex policy is “validated” when the benchmark rarely stresses it. Future BAS cohorts should register mechanism-coverage requirements (session-only, endpoint-only, hub over-merge, late horizon, beyond- W_{max} span) before claiming

Table 4: Primary M1–M7 scenario families (simplified). “Live EB1–EB4 delta?” reports whether fair builders differed on the sealed primary cohort for the intended advanced mechanism.

ID	Ordered activity (summary)	Expected EDR/NDR	Intended mechanism	Live EB1–EB4?
M1	recon → local action	process; DNS/flow	entity/time	identical
M2	remote session + SCP archive	process; SSH/SCP flows	session / file_transfer token	identical; token gap
M3	privilege / persistence steps	process; file	causal/entity	identical
M4	lateral-style dual-host steps	process; flows	session/endpoint	identical
M5	staged tooling / execution	process; file/net	entity/time	identical
M6	cleanup / cover tracks	process; file	entity/time	identical
M7	boundary actions at 0/300/3900 s	process; flows	span boundary	boundary failure

Table 5: Primary live M1–M7 cells under production formation (mean ± sample SD; $n = 5$).

Family	Mode	Coverage	Frag.	Merge	Contam.	F1
M1	attack-only	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M1	concurrent-benign	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.300 ± 0.112	1.000 ± 0.000
M2	attack-only	0.750 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M2	concurrent-benign	0.750 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.350 ± 0.091	1.000 ± 0.000
M3	attack-only	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M3	concurrent-benign	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.350 ± 0.091	1.000 ± 0.000
M4	attack-only	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M4	concurrent-benign	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.367 ± 0.217	1.000 ± 0.000
M5	attack-only	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M5	concurrent-benign	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.433 ± 0.091	1.000 ± 0.000
M6	attack-only	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	1.000 ± 0.000
M6	concurrent-benign	1.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	0.333 ± 0.257	1.000 ± 0.000
M7	attack-only	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	undef.	0.000 ± 0.000
M7	concurrent-benign	0.000 ± 0.000	0.000 ± 0.000	0.000 ± 0.000	undef.	0.000 ± 0.000

Table 6: Fair sliding-gap baselines versus legacy fixed bins on the sealed 70-run M1–M7 malicious cohort (mean ± sample SD). Primary comparisons share $W_{gap} = 300$ s and $W_{max} = 3600$ s.

Builder	Role	Coverage	Exact-comp.	Frag.	Contam.	F1
EB0_FIXED_BIN_300s	legacy	0.807 ± 0.348	0.671 ± 0.473	0.071 ± 0.259	0.178 ± 0.208	0.814 ± 0.353
EB1_SLIDING_TIME	primary	0.836 ± 0.335	0.729 ± 0.448	0.000 ± 0.000	0.175 ± 0.207	0.867 ± 0.338
EB2_SLIDING_ENTITY	primary	0.821 ± 0.349	0.714 ± 0.455	0.000 ± 0.000	0.178 ± 0.208	0.857 ± 0.352
EB3_RELATION_NO_ADVANCED_GUARDS	primary	0.836 ± 0.335	0.729 ± 0.448	0.000 ± 0.000	0.175 ± 0.207	0.863 ± 0.341
EB4_TRUSTEPIISODE	primary	0.821 ± 0.349	0.714 ± 0.455	0.000 ± 0.000	0.178 ± 0.208	0.857 ± 0.352

Table 7: Paired differences EB4–EB1_SLIDING_TIME (contamination/fragmentation flipped so positive favors EB4). Bootstrap: 2000 run-cluster resamples, seed 20260721. Zero deltas reported as no observed difference.

Metric	Mean Δ	Median Δ	95% CI	W/T/L	Interpretation
action_coverage	−0.014	0.000	[−0.043, 0.000]	0/69/1	not resolved
exact_complete_recovery	−0.014	0.000	[−0.043, 0.000]	0/69/1	not resolved
fragmentation	0.000	0.000	[0.000, 0.000]	0/70/0	no observed difference
contamination	0.000	0.000	[0.000, 0.000]	0/60/0	no observed difference
episode_f1	−0.010	0.000	[−0.029, 0.000]	0/69/1	not resolved

relation-aware gains. M7 intentionally places the final action at 3900 s, outside the frozen $W_{max} = 3600$ s span; production EB4’s zero coverage/F1 is the expected boundary-stress outcome rather than a collection failure.

4.6 Deployed Case-Materialization Trace

We additionally recorded a read-only trace from the deployed investigation surface on 2026-07-23. This trace is artifact-level integration evidence, not another experimental sample and not part of any estimator above. Figure 3(a) shows six workspace-scoped cases, five marked high risk, with both native and NDR-derived sources retained. The selected laboratory case in Figure 3(b) materializes three normalized events—one Suricata NDR event and

two Wazuh events—and exposes evidence references for ATT&CK T1018 and T1041. The interface’s 98% operational confidence and isolate_or_investigate recommendation belong to the downstream case engine; they are *not* TRUSTEPIISODE’s nullable calibrated P and are not used as evidence of model accuracy. The trace supports only the narrower claim that source-separated evidence and technique support survive into a human-reviewable case view.

4.7 Inclusive Temporal Boundary Fixtures

Unit fixtures on the production candidate function confirm inclusive equality at the frozen contracts (Table 8). For W_{gap} , 299 and 300 s are eligible while 301 s is not. For W_{max} , 3599 and 3600 s are eligible while 3601 and 3900 s are not. The 3900 s case matches the sealed



(a) Workspace-scoped case queue and selected case summary.



(b) Source-separated evidence cards and ATT&CK support map.

Figure 3: Read-only deployed case trace from the isolated TrustEpisode laboratory. The capture panels are cropped to remove account-identifying chrome. They demonstrate end-to-end evidence materialization, not calibrated probability or LLM effectiveness.

Table 8: Inclusive W_{gap}/W_{max} boundary fixtures (unit; all pass).

Dimension	Seconds	Expected	Observed
W_{gap}	299	eligible	eligible
W_{gap}	300	eligible	eligible
W_{gap}	301	ineligible	ineligible
W_{max}	3599	eligible	eligible
W_{max}	3600	eligible	eligible
W_{max}	3601	ineligible	ineligible
W_{max}	3900	ineligible	ineligible

M7 schedule offset and is *not* a substitute for live M7 formation metrics. The companion seals the complete boundary result.

4.8 Sixty-Run Determinism Matrix

On all 60 sealed malicious runs (M1–M6 $\times$ attack-only/concurrent-benign $\times$ five), D1 canonical ordered replay (10 $\times$ per run; 600 executions) is byte-identical (60/60). D2 separates raw delivery-order robustness (0/60) from re-canonicalized equality (60/60): we do *not* claim raw-order robustness. D3 ingest-delay grids (including inside, at, and beyond the 900 s revision horizon) pass 60/60; D4 duplicate/idempotent/conflict cases pass 60/60. D5 documents single-writer-per-run ownership: multi-run parallelism is not same-run concurrent mutation, and we claim neither lock-free nor multi-writer determinism.

4.9 Audit Failure Conformance (A1–A12)

RQ2 is supported by the full failure matrix: all seven Assembly-FailureRecord codes (7/7) appear with expected terminal outcomes, plus precedence, exclusivity, idempotent success retry, and append-only rejection (18/18 fixtures pass). This is synthetic forensic audit conformance, not field effectiveness.

4.10 Evidence-Loss Taxonomy Validation

The taxonomy is conformance-tested across A–G and instantiated on one sealed live case, M2 (8/8 fixture rows; 1 live). Synthetic passes do not imply field-validated coverage for every class.

4.11 M2 Evidence Waterfall

Final class: D_token_mapping_failure. Action coverage 0.75 with episode F1 1.0 is therefore not evidence that Episode Synthesis dropped an already-normalized file_transfer event. Primary sealed results are unchanged; any mapper fix requires a new version digest and a post-analysis cohort.

4.12 Contract Ablation Baselines (AB0–AB3)

AB0–AB3 are reference implementations for an isolation study: which contract property prevents which forensic failure mode. They are reference implementations for isolation, not product comparisons.

4.13 Agent-Facing Boundary Conformance

AI0 accepts an unchanged canonical AuditRecord. AI1–AI4 inject a generated fact, leak a label, force an unsupported 0.98 probability, or remove provenance; the closed schema rejects each. AI5–AI6 change a revision or evidence digest without resealing the record; canonical digest verification rejects both. AI7 changes the allowed *D* field and recomputes the unkeyed RFC 8785/SHA-256 digest; validation accepts it, as expected. Thus 8/8 cards match their specifications. The negative control shows that schema/digest validation provides integrity checks, not writer authentication. These are synthetic canonical-record checks, not LLM factuality, deployed authorization, or agent-runtime security tests.

Table 9: Determinism matrix on all 60 sealed malicious runs (D1–D5).

Test	Runs	Variants	Expected invariant	Pass	Fail	Interpretation
D1 canonical replay	60	10×/run	byte-identical digests	60	0	canonical ordered input
D2 raw delivery order	60	6 perms	equals baseline without reorder	0	60	not order-robust
D2 re-canonicalized	60	6 perms	equals baseline after sort	60	0	canonicalization required
D3 ingest delay	60	horizon grid	successor/late invariants	60	0	prior digest preserved
D4 duplicates	60	conflict/idempotent	reject or no-op	60	0	fail-closed conflicts
D5 single-writer	architecture	ownership	no same-run multi-writer	documented	0	not lock-free concurrency

Table 10: A1–A12 audit conformance fixtures (synthetic; covers all seven AssemblyFailureRecord codes).

Fixture	Expected	Observed	Term.	Prior chg?	Result
A1_SUCCESS	success	success	1	no	pass
A2_SCHEMA_INVALID	schema_invalid	schema_invalid	1	no	pass
A3_KEY_MISMATCH	key_mismatch	key_mismatch	1	no	pass
A4_CONFLICTING_DUPLICATE	conflicting_duplicate	conflicting_duplicate	1	no	pass
A5_DIGEST_MISMATCH	digest_mismatch	digest_mismatch	1	no	pass
A6_VERSION_MISMATCH	version_mismatch	version_mismatch	1	no	pass
A7_MISSING_INPUT	missing_input	missing_input	1	no	pass
A8_ASSEMBLY_TIMEOUT	assembly_timeout	assembly_timeout	1	no	pass
A9_FAILURE_PRECEDENCE_SCHEMA	schema_invalid	schema_invalid	1	no	pass
A9_FAILURE_PRECEDENCE_KEY	key_mismatch	key_mismatch	1	no	pass
A9_FAILURE_PRECEDENCE_DUPLICATE	conflicting_duplicate	conflicting_duplicate	1	no	pass
A9_FAILURE_PRECEDENCE_DIGEST	digest_mismatch	digest_mismatch	1	no	pass
A9_FAILURE_PRECEDENCE_VERSION	version_mismatch	version_mismatch	1	no	pass
A9_FAILURE_PRECEDENCE_MISSING	missing_input	missing_input	1	no	pass
A9_FAILURE_PRECEDENCE_TIMEOUT	assembly_timeout	assembly_timeout	1	no	pass
A10_TERMINAL_EXCLUSIVITY	rejected	rejected	1	no	pass
A11_IDEMPOTENT_SUCCESS_RETRY	idempotent	idempotent	1	no	pass
A12_APPEND_ONLY_STORE	both_rejected	both_rejected	1	no	pass

Table 11: A–G evidence-loss taxonomy conformance (synthetic A–G plus one sealed live M2 as D).

True class	Predicted	Source	Result
A_execution_failure	A_execution_failure	synthetic	pass
B_sensor_observability_failure	B_sensor_observability_failure	synthetic	pass
C_parser_or_normalization_failure	C_parser_or_normalization_failure	synthetic	pass
D_token_mapping_failure	D_token_mapping_failure	synthetic	pass
E_formation_membership_failure	E_formation_membership_failure	synthetic	pass
F_offline_attribution_failure	F_offline_attribution_failure	synthetic	pass
G_unknown_or_incomplete_trace	G_unknown_or_incomplete_trace	synthetic	pass
D_token_mapping_failure	D_token_mapping_failure	live M2	pass

Table 12: M2 evidence-loss waterfall (10 sealed runs).

Stage	Expected	Observed	Status
CALDERA receipt (linux.scp_archive)	10	10	retained
Raw EDR/NDR observability (scp mentions)	10	10	retained
NormalizedEvent as file_transfer	10	0	drop
Normalized as remote_session instead	10	10	remapped
Episode retains file_transfer	10	0	absent token
Episode retains remote_session	10	10	retained

4.14 Held-Out Local-Model Grounding Audits

After a separate unscored 18-card interface smoke test, we froze the system prompt, output schema, four newly valued and re-digested schema-valid AuditRecords, and 20 held-out tasks. We then ran qwen2.5:7b (model digest prefix 845dbda0ea48, Ollama 0.20.5) once with temperature 0 and seed 20260723. Tasks require exact scalar extraction, null/missing-field abstention, or resistance to a task instruction that supplies a false value. The frozen artifact retains prompts, raw responses, token counts, card/record/prompt digests, and per-field scoring.

After observing Qwen, but before downloading or querying the deployment configuration’s default gemma3:1b, we timestamped a locked replication protocol fixing the same prompt, schema, cards,

records, seed, and scoring. Gemma (digest prefix 8648f39daa8f) then ran once. Because this replication was registered after the Qwen result, Table 15 is an exploratory transfer audit, not a preregistered or statistical model comparison.

Qwen yields 15/20 exact answers, 15/20 exact citation sets, and 9/20 fully exact responses; one pointer is invalid. Gemma yields only 15/20 schema-shaped responses, 7/20 exact answers, 0/20 exact citation sets, and 0/20 fully exact responses, with 18 invalid pointers. Its raw H17, H18, and H20 responses echo injected numeric or object fragments; this lexical observation was not a frozen metric and is illustrative. The observed cross-model variability is the point of the boundary: typed input makes failure measurable, but does not make a model accurate. Neither run is field effectiveness, authorization testing, or evidence for general LLM safety.

4.15 Mechanism Conformance Cards

T1–T8 are synthetic. T2/T3 were redesigned so that specific_entity is disabled and session/endpoint becomes the sole legal relation; ablation then fragments. Passes are conformance, not live effectiveness. In the live cohort, EB4 matches EB2 but differs from EB1/EB3 on one concurrent-benign M7 run.

Table 13: Contract ablation / reference-implementation comparison (synthetic isolation study; not SOTA product baselines).

Property	AB0 Mutable	AB1 Noncan.	AB2 Silent	TrustEpisode	Evidence
prior_history_immutable_under_late	no	no	yes	yes	synthetic conformance
byte_identical_replay_under_ref_reorder	yes	no	yes	yes	synthetic conformance
explicit_terminal_failure_on_missing_input	no	no	no	yes	synthetic conformance
append_only_terminal_store	no	yes	yes	yes	synthetic conformance

Table 14: Downstream-agent canonical-record boundary cards (synthetic conformance only; not LLM accuracy or deployed authorization).

Card	Operation	Guard	Expected	Observed	Result
AI0	consume unchanged canonical AuditRecord	closed schema plus canonical digest	accept	accept	pass
AI1	add agent_asserted_fact to canonical record	additionalProperties false	schema reject	schema reject	pass
AI2	add malicious_label to online record	additionalProperties false	schema reject	schema reject	pass
AI3	replace out-of-support null P with 0.98	support-probability coupling	schema reject	schema reject	pass
AI4	remove ordered evidence references	required provenance field	schema reject	schema reject	pass
AI5	swap revision digest without resealing record	canonical digest binding	digest reject	digest reject	pass
AI6	swap evidence content digest without resealing record	canonical digest binding	digest reject	digest reject	pass
AI7	replace allowed D and recompute unkeyed digest	external writer authorization required	accept	accept	pass

Table 15: Held-out AuditRecord pilots. Gemma was locked after the Qwen result but before model download. Exact requires answer, abstention flag, and citation set to match.

Model	JSON	Answer	Abst.	Cite	Exact	Bad ptr.
Qwen2.5-7B	20/20	15/20	14/20	15/20	9/20	1
Gemma3-1B	15/20	7/20	13/20	0/20	0/20	18

4.16 Supplemental Partition (Not Primary)

Artifacts below live under a separate supplemental_mechanism partition. They do *not* enter Table 5 or the 70-run fair baseline estimator. Evidence classes are stated per row.

- **L1 (live base + offline ablation):** three sealed live M3 SSH bases. With only session linking enabled, the offline ablation links each base; disabling it fragments all three. This is not organic session-only BAS stress.
- **L2 (offline ablation on L1 seals):** on the same seals, endpoint-only linking is denser than the endpoint-disabled ablation (3/3). This is an offline mechanism check.
- **L3 (offline ablation + controlled inject):** organic hub load is *not* observed (hub-on vs hub-off identical, 0/3 organic). Controlled hub-load inject on the live sealed bases (cutoff+1 peers) makes hub-on vs hub-off diverge (3/3). Controlled inject $\neq$ organic BAS hub stress.
- **L4 late evidence (two distinct classes):** (i) *controlled late inject* on live seals yields successors and beyond_revision_horizon records (3/3). (ii) A *wall-clock hold/release* uses real delays: within-horizon 180 s holds seal and show multi-revision successors (2/2 pass); beyond-horizon attempts (930 s hold-release; 1080 s quiet wait) *seal* but do *not* emit LateEvidenceRecords under organic telemetry (2 sealed, 0/2 LateEvidence; new membership formed instead). Beyond-horizon localization is therefore supported by controlled inject, not by wall-clock beyond runs in this lab.

4.17 Artifact Reproducibility

The anonymous companion archive fails closed unless Python imports its packaged reference runtime, executes all 38 implementation tests, and re-runs A1–A12, A–G, AB0–AB3, AI0–AI7, T1–T8, and the temporal boundary fixtures. It regenerates tables from sealed JSON and checks the reported D1–D5, fair-baseline, and two local-model summaries; raw/derived equality is checked for both model audits. It does not re-run live collection, cohort re-synthesis that requires the omitted raw lab bundles, or Ollama. Every clean-room command executes inside the extracted ZIP; the adjacent report records the results. Digests are refreshed with each packing.

Ethical scope. All traffic and workloads were confined to isolated, author-controlled laboratory systems. The study used no production users, personal data, human subjects, or third-party targets; screenshots were cropped to exclude account-identifying chrome.

4.18 Limitations

Two Ubuntu targets; non-commercial EDR; $n = 5$ per cell; controlled benign workloads; standalone benign controls are not false-positive trials; no production base rates; fair sliding results largely overlap, but EB4 trails EB1/EB3 on one concurrent-benign M7 run; synthetic cards $\neq$ field validation; raw order not robust; no fitted scorer/calibration; no partial-telemetry/outage; M2 mapping not corrected in the primary cohort. Organic hub stress remains unobserved on live streams. Wall-clock *within*-horizon holds are executed (2/2); wall-clock *beyond*-horizon LateEvidenceRecord under organic telemetry was attempted (2 sealed) but not observed. L3 hub divergence and L4 beyond-horizon localization use controlled inject on live seals where noted. M7 is primary boundary stress and should not be interpreted as successful 3900 s recovery under a 3600 s span. AI0–AI7 separate record validation from writer authorization: AI7 confirms that a schema-valid reseal is accepted, while deployed service-identity enforcement remains unevaluated. The LLM audits are synthetic, two-model, same-card, and once-run; the Gemma replication was locked only after Qwen was observed. Their 9/20 and 0/20 exact rates cannot establish general factuality, a size effect, robustness, or operational utility.

Table 16: Targeted synthetic mechanism cards (conformance only; not mixed into primary effectiveness).

Card	Result	Expected	Observed
T1-PROCESS-CAUSAL	pass	fragment without causal	full-eps:1, abl-eps:2
T2-AUTH-SESSION	pass	session only links ablation fragm...	session-only:1, abl:2
T3-ENDPOINT-MAPPING	pass	endpoint only links ablation frag...	endpoint-only:1, abl:2
T4-HUB-GUARD	pass	hub guard changes membership or e...	full-eps:4, no-hub-eps:1
T5-PARENT-MERGE	pass	merge enabled creates parent epis...	full-parents:3, abl-parents:0
T6-LATE-SUCCESSOR	pass	disabling late successor increase...	full-rev:4, abl-rev:3, full-late:...
T7-BEYOND-HORIZON	pass	beyond horizon emits LateEvidence...	late:1, reasons:(beyond-revision-...
T8-TIE-BREAK	pass	canonical arrival order yields id...	equal:True; raw-reverse-differs:True

Analysis status. Complete (primary): 110 sealed live executions (70 malicious M1–M7; 40 standalone B1–B8 benign controls); fair sliding baselines with paired CI; T1–T8; M2 waterfall; 60-run determinism D1–D5; A1–A12; A–G; AB0–AB3; AI0–AI7; held-out Qwen audit + locked Gemma replication; inclusive boundary fixtures including 3900 s. Complete (supplemental, not primary): L1 live bases + offline L1/L2; L3 controlled hub inject; L4 controlled late inject; wall-clock within 2/2 and beyond 2 sealed/0 LateEvidence. Not executed: organic hub-stressed BAS traffic; fitted scorer / RP / PO1.

5 Conclusion

TRUSTEPISODE is a deterministic, auditable EDR/NDR reconstruction contract and a read-only evidence boundary for AI-assisted investigation: single-owner immutable revisions, append-only terminal outcomes with a complete failure-code matrix, and stage-localized A–G evidence-loss diagnosis. On the sealed M1–M7 cohort, sliding sessionization improves completeness relative to fixed bins, but fair sliding time/entity/relation builders do not establish TrustEpisode superiority. EB4 matches EB2 on all 70 action-coverage pairs, but trails EB1/EB3 on one concurrent-benign M7 run; the EB4–EB1 action-coverage estimate is $\Delta = -0.014$ with 95% CI $[-0.043, 0.000]$. That mostly-overlapping result is forensic evidence about benchmark coverage and one boundary-case cost: advanced guards require targeted conformance cards, which we provide without treating them as live effectiveness. Canonical replay is byte-identical on all 60 sealed malicious runs; raw arrival-order robustness requires re-canonicalization. M2’s missing file_transfer token is D_token_mapping_failure, not a demonstrated formation drop. Contract ablation shows which properties prevent mutable history, silent missing-input failure, and noncanonical replay. Eight downstream-boundary cards accept the valid canonical control, reject six schema-invalid or unsealed mutations, and accept a resealed allowed-field negative control. The latter shows why an unkeyed digest is not writer authentication; deployed authorization remains unevaluated. A separate once-run held-out Qwen2.5-7B pilot yields 15/20 exact answers but only 9/20 fully exact answer/abstention/citation triples. A same-card replication, locked after that result but before querying deployment-default Gemma3-1B, yields 7/20 exact answers, no exact citation sets, no fully exact responses, and 18 invalid pointers. This is evidence that a typed record makes model-dependent failure inspectable, not that either model is generally accurate or that the observed gap estimates a size effect. Supplemental partition artifacts (not primary estimators) include L1 live bases with offline L1/L2 ablations, L3 controlled

hub inject, L4 controlled late inject, wall-clock within-horizon holds (2/2), and sealed beyond-horizon attempts without organic LateEvidenceRecord. Sealed M7 is a primary 3900 s boundary-stress family with expected zero recovery beyond W_{max} . Fitted scoring, calibration, and telemetry-outage robustness remain unevaluated and outside these claims. A deployed read-only trace confirms that source-separated Suricata/Wazuh evidence and ATT&CK references reach the analyst case view; its displayed confidence is not treated as the contract’s probability or as LLM-effectiveness evidence. The result is a falsifiable evidence substrate for future AI reasoning; both negative audits reinforce rather than remove the need to distrust the model.

A Frozen Submission Contract

Machine and audit contract. The normative companion uses closed JSON Schema Draft 2020-12 objects [22] for NormalizedEvent, SourceCoverage, EpisodeRevision, LateEvidenceRecord, FeatureVector, SupportDecision, AuditRecord, and AssemblyFailureRecord. FeatureVector cannot consume health state. SupportDecision resolves detector availability, support, group key, and eligibility; zero eligible detectors forces out-of-support, false eligibility, and JSON null probability. For each revision, the Assembler inserts exactly one terminal AuditRecord after a complete version-matched join or one AssemblyFailureRecord after schema, key, duplicate, digest, version, missing-input, or timeout failure. Failure cannot later be replaced by success. Online objects reject labels, scenario/card IDs, receipts, and attack windows. Canonical bytes follow RFC 8785 [13]; application-ordered reference arrays are sorted before SHA-256, and every schema, parser, formation, detector, model, support, sufficiency, assembly, and manifest version binds both identifier and content digest.

Episode and lifecycle constants. Frozen values are $W_{gap} = 300$ s, $W_{max} = 3600$ s, watermark lag 120 s, revision horizon 900 s, and closed-lineage retention 86400 s. Candidate relations rank causal/session/endpoint/ entity as (4, 3, 2, 1) and sort by relation rank, number of satisfied classes, time distance, start time, then episode ID. A hub has at least 25 relations in 600 s; merge admits at most three causally/session/endpoint-connected parents and 500 events/200 entities. Events sort lexicographically by event time, ingest time, source-instance bytes, and event-ID bytes. First finalization fixes the revision deadline. A qualifying event before it creates an immutable successor without extending the horizon; a later related event creates only a LateEvidenceRecord. Replay equality covers states, successors, late records, ordered references, and digests.

Evaluation and artifact boundary. The achieved primary cohort contains five valid runs per M1–M7 card/mode and five standalone runs per B1–B8 card. The B1–B8 runs are execution controls, not false-positive trials because no fitted decision rule is evaluated. Offline matching freezes

malicious coverage 0.50, contamination 0.25, match weight 0.50, temporal IoU 0.10, entity Jaccard 0.50, and a 2000-resample run-cluster bootstrap with seed 20260721. Fitted calibration, RP0-RP7 partial-telemetry replay, physical outage, and full cost benchmarking remain unexecuted and are not claimed. The companion contains the full schemas, thresholds, cards, validators, sealed result copies, and claim matrix. Its filename is TrustEpisode_reproducibility_companion_v1.zip; concatenating the following lines gives its SHA-256:

```
cc4d451b50f975fdd7ab07c263a8ce20
44ab771126808e485faac4033852a730
```

References

- [1] Frank Breitering, Hudan Studiawan, and Chris Hargreaves. 2025. SoK: Timeline based event reconstruction for digital forensics: Terminology, methodology, and current challenges. *Forensic Science International: Digital Investigation* 53 (2025), 301932. DFRWS USA 2025. doi:10.1016/j.fsidi.2025.301932
- [2] Joshua Bruntz. 2023. Validation of forensic tools and methods: A primer for the digital forensics examiner. *Wiley Interdisciplinary Reviews: Forensic Science* 5, 2 (2023), e1474. doi:10.1002/wfs2.1474
- [3] Brian Carrier. 2005. *File System Forensic Analysis*. Addison-Wesley, Boston, MA, USA.
- [4] Eoghan Casey. 2011. *Digital Evidence and Computer Crime: Forensic Science, Computers, and the Internet* (3 ed.). Academic Press, Waltham, MA, USA.
- [5] Haodong Chen, Ziheng Zhang, Jinghui Jiang, Qiang Su, and Qiao Xiang. 2026. Holmes: An Evidence-Grounded LLM Agent for Auditable DDoS Investigation in Cloud Networks. arXiv:2601.14601 [cs.CR] doi:10.48550/arXiv.2601.14601
- [6] Zijun Cheng, Qiujian Lv, Jinyuan Liang, Yan Wang, Degang Sun, Thomas Pasquier, and Xueyuan Han. 2024. KAIROS: Practical Intrusion Detection and Investigation using Whole-system Provenance. In *2024 IEEE Symposium on Security and Privacy (SP)*. IEEE, Los Alamitos, CA, USA, 3533–3551. doi:10.1109/SP54263.2024.00005
- [7] Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. Enabling Large Language Models to Generate Text with Citations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Singapore, 6465–6488. doi:10.18653/v1/2023.emnlp-main.398
- [8] Xueyuan Han, Thomas F. J. Pasquier, Adam Bates, James Mickens, and Margo I. Seltzer. 2020. UNICORN: Runtime Provenance-Based Detector for Advanced Persistent Threats. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, San Diego, CA, USA, 18 pages. doi:10.14722/ndss.2020.24046
- [9] Shaofei Li, Feng Dong, Xusheng Xiao, Haoyu Wang, Fei Shao, Jiedong Chen, Yao Guo, Xiangqun Chen, and Ding Li. 2024. NODLINK: An Online System for Fine-Grained APT Attack Detection and Investigation. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, Reston, VA, USA, 18 pages. doi:10.14722/ndss.2024.23204
- [10] Sadeq M. Milajerdi, Rigel Gjomemo, Birhanu Eshete, R. Sekar, and V. N. Venkatakrishnan. 2019. HOLMES: Real-time APT Detection through Correlation of Suspicious Information Flows. In *2019 IEEE Symposium on Security and Privacy (SP)*. IEEE, Los Alamitos, CA, USA, 1137–1152. doi:10.1109/SP.2019.00026
- [11] National Institute of Standards and Technology. 2006. *Guide to Integrating Forensic Techniques into Incident Response*. Technical Report NIST Special Publication 800-86. U.S. Department of Commerce. doi:10.6028/NIST.SP.800-86
- [12] Peng Ning, Yun Cui, and Douglas S. Reeves. 2002. Constructing Attack Scenarios through Correlation of Intrusion Alerts. In *Proceedings of the 9th ACM Conference on Computer and Communications Security*. Association for Computing Machinery, New York, NY, USA, 245–254. doi:10.1145/586110.586144
- [13] Anders Rundgren, Benjamin Jordan, and Samuel Erdtman. 2020. *JSON Canonicalization Scheme (JCS)*. RFC 8785. RFC Editor. doi:10.17487/RFC8785
- [14] Scientific Working Group on Digital Evidence. 2024. *Minimum Requirements for Testing Tools Used in Digital and Multimedia Forensics*. Technical Report 18-Q-001-2.1. SWGDE. <https://www.swgde.org/18-q-001/>
- [15] Shahroz Tariq, Mohan Baruwat Chhetri, Surya Nepal, and Cecile Paris. 2025. Alert Fatigue in Security Operations Centres: Research Challenges and Opportunities. *Comput. Surveys* 57, 9 (2025), 224:1–224:38. doi:10.1145/3723158
- [16] Céline Vanini, Christopher J. Hargreaves, Harm van Beek, and Frank Breitering. 2024. Was the clock correct? Exploring timestamp interpretation through time anchors for digital forensic event reconstruction. *Forensic Science International: Digital Investigation* 49 (2024), 301759. DFRWS USA 2024. doi:10.1016/j.fsidi.2024.301759
- [17] Mehrnoush Vaseghpanah, Sam Jabbehdari, and Hamidreza Navidi. 2026. Auditing Inferential Blind Spots: A Framework for Evaluating Forensic Coverage in Network Telemetry Architectures. *Network* 6, 1 (2026), 9. doi:10.3390/network6010009

- [18] Gregory S. Wales. 2026. A research-focused framework for empirical method validation in digital and multimedia evidence. *Journal of Forensic Sciences* 71, 2 (2026), 771–784. doi:10.1111/1556-4029.70253
- [19] Jonas Wallat, Maria Heuss, Maarten de Rijke, and Avishek Anand. 2025. Correctness is not Faithfulness in Retrieval Augmented Generation Attributions. In *Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval*. Association for Computing Machinery, New York, NY, USA, 22–32. doi:10.1145/3731120.3744592
- [20] Lingzhi Wang, Xiangmin Shen, Weijian Li, Zhenyuan Li, R. Sekar, Han Liu, and Yan Chen. 2025. Incorporating Gradients to Rules: Towards Lightweight, Adaptive Provenance-based Intrusion Detection. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, Reston, VA, USA, 18 pages. doi:10.14722/ndss.2025.230822
- [21] Bowen Wei, Yuan Shen Tay, Howard Liu, Jinhao Pan, Kun Luo, Ziwei Zhu, and Chris Jordan. 2025. CORTEX: Collaborative LLM Agents for High-Stakes Alert Triage. arXiv:2510.00311 [cs.CL] doi:10.48550/arXiv.2510.00311
- [22] Austin Wright, Henry Andrews, Ben Hutton, and Greg Dennis. 2022. JSON Schema: A Media Type for Describing JSON Documents, Draft 2020-12. JSON Schema specification. Published 2022-06-16. <https://json-schema.org/draft/2020-12/json-schema-core.html>
- [23] Renjie Wu, Hu Wang, Hsiang-Ting Chen, and Gustavo Carneiro. 2024. Deep Multimodal Learning with Missing Modality: A Survey. arXiv:2409.07825 [cs.CV] doi:10.48550/arXiv.2409.07825
- [24] Yuhao Wu, Franziska Roesner, Tadayoshi Kohno, Ning Zhang, and Umar Iqbal. 2025. IsolateGPT: An Execution Isolation Architecture for LLM-Based Agentic Systems. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, San Diego, CA, USA, 20 pages. doi:10.14722/ndss.2025.241131
- [25] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *The Eleventh International Conference on Learning Representations*. OpenReview.net, Kigali, Rwanda, 33 pages. https://openreview.net/forum?id=WE_vluYUL-X

Generative AI Use Disclosure

OpenAI Codex and Cursor agents assisted with language editing, experiment orchestration, consistency checking, figure preparation, citation verification, and claim-to-evidence audits. The authors independently designed the study, verified the code, experimental outputs, citations, screenshots, and reported results, and take responsibility for the manuscript.